

```
In [ ]: # Reference: Thanks to
# Moon, T. (2018, November 23). Tjmoon0104/pytorch-tiny-imagenet. Retrieved June 14,
# 2020, from https://github.com/tjmoon0104/pytorch-tiny-imagenet

#Zakka, K. (2017). Train, Validation and Test Split for torchvision Data
sets. Retrieved June
#14, 2020, from https://gist.github.com/kevinzakka/d33bf8d6c7f06a9d8
c76d97a7879f5cb

#Xiangjianxiaolu(2018, September 25). Pytorch入门实战: ResNet18图像分类 (Cifar10) .
#Retrieved June 14, 2020, from http://flyrie.top/2018/09/26/Pytorch_
ResNet18_Cifar10/
```

setup

```
In [2]: pip install livelossplot
```

Requirement already satisfied: livelossplot in ./local/lib/python3.7/site-packages (0.5.1)

Requirement already satisfied: bokeh; python_version >= "3.6" in /opt/conda/lib/python3.7/site-packages (from livelossplot) (1.3.4)

Requirement already satisfied: matplotlib; python_version >= "3.6" in /opt/conda/lib/python3.7/site-packages (from livelossplot) (3.1.1)

Requirement already satisfied: ipython in /opt/conda/lib/python3.7/site-packages (from livelossplot) (7.7.0)

Requirement already satisfied: six>=1.5.2 in /opt/conda/lib/python3.7/site-packages (from bokeh; python_version >= "3.6"->livelossplot) (1.12.0)

Requirement already satisfied: tornado>=4.3 in /opt/conda/lib/python3.7/site-packages (from bokeh; python_version >= "3.6"->livelossplot) (6.0.3)

Requirement already satisfied: python-dateutil>=2.1 in /opt/conda/lib/python3.7/site-packages (from bokeh; python_version >= "3.6"->livelossplot) (2.8.0)

Requirement already satisfied: numpy>=1.7.1 in /opt/conda/lib/python3.7/site-packages (from bokeh; python_version >= "3.6"->livelossplot) (1.16.4)

Requirement already satisfied: PyYAML>=3.10 in /opt/conda/lib/python3.7/site-packages (from bokeh; python_version >= "3.6"->livelossplot) (5.1.2)

Requirement already satisfied: packaging>=16.8 in /opt/conda/lib/python3.7/site-packages (from bokeh; python_version >= "3.6"->livelossplot) (19.0)

Requirement already satisfied: pillow>=4.0 in /opt/conda/lib/python3.7/site-packages (from bokeh; python_version >= "3.6"->livelossplot) (6.1.0)

Requirement already satisfied: Jinja2>=2.7 in /opt/conda/lib/python3.7/site-packages (from bokeh; python_version >= "3.6"->livelossplot) (2.10.1)

Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /opt/conda/lib/python3.7/site-packages (from matplotlib; python_version >= "3.6"->livelossplot) (2.4.2)

Requirement already satisfied: kiwisolver>=1.0.1 in /opt/conda/lib/python3.7/site-packages (from matplotlib; python_version >= "3.6"->livelossplot) (1.1.0)

Requirement already satisfied: cyclor>=0.10 in /opt/conda/lib/python3.7/site-packages (from matplotlib; python_version >= "3.6"->livelossplot) (0.10.0)

Requirement already satisfied: decorator in /opt/conda/lib/python3.7/site-packages (from ipython->livelossplot) (4.4.0)

Requirement already satisfied: pexpect; sys_platform != "win32" in /opt/conda/lib/python3.7/site-packages (from ipython->livelossplot) (4.7.0)

Requirement already satisfied: jedi>=0.10 in /opt/conda/lib/python3.7/site-packages (from ipython->livelossplot) (0.14.1)

Requirement already satisfied: backcall in /opt/conda/lib/python3.7/site-packages (from ipython->livelossplot) (0.1.0)

Requirement already satisfied: setuptools>=18.5 in /opt/conda/lib/python3.7/site-packages (from ipython->livelossplot) (41.0.1)

Requirement already satisfied: prompt-toolkit<2.1.0,>=2.0.0 in /opt/conda/lib/python3.7/site-packages (from ipython->livelossplot) (2.0.9)

Requirement already satisfied: traitlets>=4.2 in /opt/conda/lib/python3.7/site-packages (from ipython->livelossplot) (4.3.2)

Requirement already satisfied: pickleshare in /opt/conda/lib/python3.7/site-packages (from ipython->livelossplot) (0.7.5)
Requirement already satisfied: pygments in /opt/conda/lib/python3.7/site-packages (from ipython->livelossplot) (2.4.2)
Requirement already satisfied: MarkupSafe>=0.23 in /opt/conda/lib/python3.7/site-packages (from Jinja2>=2.7->bokeh; python_version >= "3.6"->livelossplot) (1.1.1)
Requirement already satisfied: ptyprocess>=0.5 in /opt/conda/lib/python3.7/site-packages (from pexpect; sys_platform != "win32"->ipython->livelossplot) (0.6.0)
Requirement already satisfied: parso>=0.5.0 in /opt/conda/lib/python3.7/site-packages (from jedi>=0.10->ipython->livelossplot) (0.5.1)
Requirement already satisfied: wcwidth in /opt/conda/lib/python3.7/site-packages (from prompt-toolkit<2.1.0,>=2.0.0->ipython->livelossplot) (0.1.7)
Requirement already satisfied: ipython_genutils in /opt/conda/lib/python3.7/site-packages (from traitlets>=4.2->ipython->livelossplot) (0.2.0)
Note: you may need to restart the kernel to use updated packages.

In [3]: `pip install --user opencv-python`

Requirement already satisfied: opencv-python in ./local/lib/python3.7/site-packages (4.2.0.34)
Requirement already satisfied: numpy>=1.14.5 in /opt/conda/lib/python3.7/site-packages (from opencv-python) (1.16.4)
Note: you may need to restart the kernel to use updated packages.

In [4]: `import torch, os
import torch.nn as nn
import torch.optim as optim
import torchvision.datasets as datasets
import torch.utils.data as data
import torchvision.transforms as transforms
import torchvision.models as models
from train_model import train_model
from test_model import test_model
%matplotlib inline
import torchvision
import torch.utils.model_zoo as model_zoo
from torch.utils.data.sampler import SubsetRandomSampler
from data_loader import get_train_valid_loader
from data_loader import get_test_loader
import torch.nn.functional as F`

data loader

```

In [5]: transform = transforms.Compose(
        [transforms.ToTensor(),
         transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

trainloader, valloader = get_train_valid_loader('./data', batch_size=10,
augment=0, random_seed=34567,
                                                valid_size=0.02, shuffle
= True, num_workers=2, pin_memory=True)

testset = torchvision.datasets.CIFAR10(root='./data', train=False, downl
oad=True, transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=4, shuffle=
False, num_workers=2)

dataloaders = {'train': trainloader, 'val': valloader, 'test': testloade
r}
dataset_sizes = {'train': 49000, 'val': 1000, 'test': len(testset)}

classes = ('plane', 'car', 'bird', 'cat',
           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')

```

Files already downloaded and verified
Files already downloaded and verified
Files already downloaded and verified

RESNET

```

In [6]: class ResidualBlock(nn.Module):
    def __init__(self, in_channels, out_channels, stride=1, **kwargs):
        super().__init__(**kwargs)
        self.residual = nn.Sequential(
            nn.Conv2d(in_channels=in_channels, out_channels=out_channels
, kernel_size=3, stride=stride, padding=1, bias=False),
            nn.BatchNorm2d(out_channels),
            nn.ReLU(),
            nn.Conv2d(in_channels=out_channels, out_channels=out_channels
s, kernel_size=3, stride=1, padding=1, bias=False),
            nn.BatchNorm2d(out_channels),
        )

        self.identity = nn.Sequential()
        if stride != 1 or in_channels != out_channels:
            self.identity = nn.Sequential(
                nn.Conv2d(in_channels=in_channels, out_channels=out_channels
nels, kernel_size=1, stride=stride)
            )

    def forward(self, x):
        out = self.residual(x)
        out += self.identity(x)
        out = F.relu(out)
        return out

class ResNet(nn.Module):
    def __init__(self, num_class=10):
        super().__init__()
        self.conv1 = nn.Sequential(
            nn.Conv2d(3, 64, 3, 1, 1),
            nn.BatchNorm2d(64),
            nn.ReLU(),
        )
        self.layer1 = self._make_layer(64, 64, 2)
        self.layer2 = self._make_layer(64, 128, 2, 2)
        self.layer3 = self._make_layer(128, 256, 2, 2)
        self.layer4 = self._make_layer(256, 512, 2, 2)
        self.avg_pool = nn.AvgPool2d(4)
        self.linear = nn.Linear(512, num_class)

    def _make_layer(self, in_channel, out_channel, bloch_num, stride=1):
        blocks = []
        blocks.append(ResidualBlock(in_channel, out_channel, stride))
        for i in range(1, bloch_num):
            blocks.append(ResidualBlock(out_channel, out_channel))
        return nn.Sequential(*blocks)

    def forward(self, x):
        x = self.conv1(x)
        x = self.layer1(x)
        x = self.layer2(x)
        x = self.layer3(x)
        x = self.layer4(x)
        x = self.avg_pool(x)

```


```
x = x.view(x.shape[0], -1)
return x
```

```
In [15]: def resnet(pretrained=False, **kwargs):
          r"""AlexNet model architecture from the
          `One weird trick...' <https://arxiv.org/abs/1404.5997>`_ paper.
          Args:
              pretrained (bool): If True, returns a model pre-trained on Image
Net
          """
          model = ResNet(**kwargs)
          # if pretrained:
          #     model.load_state_dict(model_zoo.load_url(model_urls['alexnet']))
          #     model.classifier[1] = nn.Linear(256 * 2 * 2, 4096)
          #     model.classifier[6] = nn.Linear(4096, 10)
          return model
```

Train

```
In [10]: #Load Resnet18
model_ft = resnet()
#Finetune Final few layers to adjust for tiny imagenet input
# model_ft.avgpool = nn.AdaptiveAvgPool2d(1)
# model_ft.fc.out_features = 10
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
model_ft = model_ft.to(device)
#Loss Function
criterion = nn.CrossEntropyLoss()
# Observe that all parameters are being optimized
optimizer_ft = optim.SGD(model_ft.parameters(), lr=0.001, momentum=0.9)

#Train
train_model("64_no_pre_kelly", model_ft, dataloaders, dataset_sizes, criterion, optimizer_ft, num_epochs=10)
```



```
accuracy
    training          (min:    0.476, max:    0.893, cur:
0.893)
    validation        (min:    0.654, max:    0.830, cur:
0.830)
log loss
    training          (min:    0.306, max:    1.437, cur:
0.306)
    validation        (min:    0.523, max:    0.962, cur:
0.523)
Train Loss: 0.3057 Acc: 0.8930
Val Loss: 0.5232 Acc: 0.8300
```

```
Training complete in 34m 23s
Best Validation Accuracy: 0.8300000000000001, Epoch: 10
```

Test


```
In [25]: #Test Resnet18
model_ft = resnet()
#Finetune Final few layers to adjust for tiny imagenet input
# model_ft.avgpool = nn.AdaptiveAvgPool2d(1)
# model_ft.fc.out_features = 10
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
model_ft.load_state_dict(torch.load('models/64_no_pre_kelly/model_10_epoch.pt')) #TODO
model_ft = model_ft.to(device)

#Loss Function
criterion = nn.CrossEntropyLoss()
# Observe that all parameters are being optimized
optimizer_ft = optim.SGD(model_ft.parameters(), lr=0.001, momentum=0.9)

#Test
test_model(model_ft, dataloaders, dataset_sizes, criterion, optimizer_ft
)
```

Iteration: 2500/2500, Loss: 5.67810583114624....

Test Loss: 0.9124 Acc: 0.7029

Test complete in 0m 27s

RESNET -activation: leakyrelu

```

In [30]: class ResidualBlock(nn.Module):
    def __init__(self, in_channels, out_channels, stride=1, **kwargs):
        super().__init__(**kwargs)
        self.residual = nn.Sequential(
            nn.Conv2d(in_channels=in_channels, out_channels=out_channels
, kernel_size=3, stride=stride, padding=1, bias=False),
            nn.BatchNorm2d(out_channels),
            nn.LeakyReLU(inplace = True),
            nn.Conv2d(in_channels=out_channels, out_channels=out_channel
s, kernel_size=3, stride=1, padding=1, bias=False),
            nn.BatchNorm2d(out_channels),
        )

        self.identity = nn.Sequential()
        if stride != 1 or in_channels != out_channels:
            self.identity = nn.Sequential(
                nn.Conv2d(in_channels=in_channels, out_channels=out_chan
nels, kernel_size=1, stride=stride)
            )

    def forward(self, x):
        out = self.residual(x)
        out += self.identity(x)
        out = F.relu(out)
        return out

class ResNet(nn.Module):
    def __init__(self, num_class=10):
        super().__init__()
        self.conv1 = nn.Sequential(
            nn.Conv2d(3, 64, 3, 1, 1),
            nn.BatchNorm2d(64),
            nn.LeakyReLU(inplace = True),
        )
        self.layer1 = self._make_layer(64, 64, 2)
        self.layer2 = self._make_layer(64, 128, 2, 2)
        self.layer3 = self._make_layer(128, 256, 2, 2)
        self.layer4 = self._make_layer(256, 512, 2, 2)
        self.avg_pool = nn.AvgPool2d(4)
        self.linear = nn.Linear(512, num_class)

    def _make_layer(self, in_channel, out_channel, bloch_num, stride=1):
        blocks = []
        blocks.append(ResidualBlock(in_channel, out_channel, stride))
        for i in range(1, bloch_num):
            blocks.append(ResidualBlock(out_channel, out_channel))
        return nn.Sequential(*blocks)


    def forward(self, x):
        x = self.conv1(x)
        x = self.layer1(x)
        x = self.layer2(x)
        x = self.layer3(x)
        x = self.layer4(x)
        x = self.avg_pool(x)

```

```
x = x.view(x.shape[0], -1)
return x
```

```
In [31]: #Load Resnet18
model_ft = resnet()
#Finetune Final few layers to adjust for tiny imagenet input
# model_ft.avgpool = nn.AdaptiveAvgPool2d(1)
# model_ft.fc.out_features = 10
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
model_ft = model_ft.to(device)
#Loss Function
criterion = nn.CrossEntropyLoss()
# Observe that all parameters are being optimized
optimizer_ft = optim.SGD(model_ft.parameters(), lr=0.001, momentum=0.9)

#Train
train_model("64_no_pre_kelly_leaky", model_ft, dataloaders, dataset_size
s, criterion, optimizer_ft, num_epochs=10)
```



```
accuracy
  training (min: 0.491, max: 0.894, cur: 0.894)
  validation (min: 0.648, max: 0.826, cur: 0.825)
log loss
  training (min: 0.304, max: 1.411, cur: 0.304)
  validation (min: 0.513, max: 0.985, cur: 0.573)
Train Loss: 0.3040 Acc: 0.8943
Val Loss: 0.5725 Acc: 0.8250
```

```
Training complete in 33m 28s
Best Validation Accuracy: 0.8260000000000001, Epoch: 7
```

```
In [32]: #Test Resnet18
model_ft = resnet()
#Finetune Final few layers to adjust for tiny imagenet input
# model_ft.avgpool = nn.AdaptiveAvgPool2d(1)
# model_ft.fc.out_features = 10
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
model_ft.load_state_dict(torch.load('models/64_no_pre_kelly_leaky/model_7_epoch.pt')) #TODO
model_ft = model_ft.to(device)

#Loss Function
criterion = nn.CrossEntropyLoss()
# Observe that all parameters are being optimized
optimizer_ft = optim.SGD(model_ft.parameters(), lr=0.001, momentum=0.9)

#Test
test_model(model_ft, dataloaders, dataset_sizes, criterion, optimizer_ft
)
```

Iteration: 2500/2500, Loss: 3.1883418560028076..

Test Loss: 0.8839 Acc: 0.7106

Test complete in 0m 27s

RESNET - optimizer: Adam

```

In [7]: class ResidualBlock(nn.Module):
    def __init__(self, in_channels, out_channels, stride=1, **kwargs):
        super().__init__(**kwargs)
        self.residual = nn.Sequential(
            nn.Conv2d(in_channels=in_channels, out_channels=out_channels
, kernel_size=3, stride=stride, padding=1, bias=False),
            nn.BatchNorm2d(out_channels),
            nn.ReLU(),
            nn.Conv2d(in_channels=out_channels, out_channels=out_channels
s, kernel_size=3, stride=1, padding=1, bias=False),
            nn.BatchNorm2d(out_channels),
        )

        self.identity = nn.Sequential()
        if stride != 1 or in_channels != out_channels:
            self.identity = nn.Sequential(
                nn.Conv2d(in_channels=in_channels, out_channels=out_channels
nels, kernel_size=1, stride=stride)
            )

    def forward(self, x):
        out = self.residual(x)
        out += self.identity(x)
        out = F.relu(out)
        return out

class ResNet(nn.Module):
    def __init__(self, num_class=10):
        super().__init__()
        self.conv1 = nn.Sequential(
            nn.Conv2d(3, 64, 3, 1, 1),
            nn.BatchNorm2d(64),
            nn.ReLU(),
        )
        self.layer1 = self._make_layer(64, 64, 2)
        self.layer2 = self._make_layer(64, 128, 2, 2)
        self.layer3 = self._make_layer(128, 256, 2, 2)
        self.layer4 = self._make_layer(256, 512, 2, 2)
        self.avg_pool = nn.AvgPool2d(4)
        self.linear = nn.Linear(512, num_class)

    def _make_layer(self, in_channel, out_channel, bloch_num, stride=1):
        blocks = []
        blocks.append(ResidualBlock(in_channel, out_channel, stride))
        for i in range(1, bloch_num):
            blocks.append(ResidualBlock(out_channel, out_channel))
        return nn.Sequential(*blocks)

    def forward(self, x):
        x = self.conv1(x)
        x = self.layer1(x)
        x = self.layer2(x)
        x = self.layer3(x)
        x = self.layer4(x)
        x = self.avg_pool(x)

```


```
x = x.view(x.shape[0], -1)
return x
```

```
In [34]: #Load AlexNet
model_ft = resnet()

device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
model_ft = model_ft.to(device)

#Loss Function
criterion = nn.CrossEntropyLoss()
# Observe that all parameters are being optimized
optimizer_ft = optim.Adam(model_ft.parameters(), lr=0.001)
```

```
In [35]: train_model("resnet_adam",model_ft, dataloaders, dataset_sizes, criterio
n, optimizer_ft, num_epochs=10)
```



```
accuracy
    training (min: 0.490, max: 0.918, cur:
0.918)
    validation (min: 0.644, max: 0.838, cur:
0.838)
log loss
    training (min: 0.235, max: 1.408, cur:
0.235)
    validation (min: 0.477, max: 0.958, cur:
0.481)
Train Loss: 0.2347 Acc: 0.9184
Val Loss: 0.4809 Acc: 0.8380
```

```
Training complete in 53m 47s
Best Validation Accuracy: 0.838, Epoch: 10
```

```
In [36]: #Test Resnet18
model_ft = resnet()
#Finetune Final few layers to adjust for tiny imagenet input
# model_ft.avgpool = nn.AdaptiveAvgPool2d(1)
# model_ft.fc.out_features = 10
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
model_ft.load_state_dict(torch.load('models/resnet_adam/model_10_epoch.p
t')) #TODO
model_ft = model_ft.to(device)

#Loss Function
criterion = nn.CrossEntropyLoss()
# Observe that all parameters are being optimized
optimizer_ft = optim.Adam(model_ft.parameters(), lr=0.001)

#Test
test_model(model_ft, dataloaders, dataset_sizes, criterion, optimizer_ft
)
```

Iteration: 2500/2500, Loss: 7.171634674072266.....

Test Loss: 0.7421 Acc: 0.7644

Test complete in 0m 28s

RESNET - fine tune

```

In [8]: class ResidualBlock(nn.Module):
    def __init__(self, in_channels, out_channels, stride=1, **kwargs):
        super().__init__(**kwargs)
        self.residual = nn.Sequential(
            nn.Conv2d(in_channels=in_channels, out_channels=out_channels
, kernel_size=3, stride=stride, padding=1, bias=False),
            nn.BatchNorm2d(out_channels),
            nn.ReLU(),
            nn.Conv2d(in_channels=out_channels, out_channels=out_channels
s, kernel_size=3, stride=1, padding=1, bias=False),
            nn.BatchNorm2d(out_channels),
        )

        self.identity = nn.Sequential()
        if stride != 1 or in_channels != out_channels:
            self.identity = nn.Sequential(
                nn.Conv2d(in_channels=in_channels, out_channels=out_channels
nels, kernel_size=1, stride=stride)
            )

    def forward(self, x):
        out = self.residual(x)
        out += self.identity(x)
        out = F.relu(out)
        return out

class ResNet(nn.Module):
    def __init__(self, num_class=10):
        super().__init__()
        self.conv1 = nn.Sequential(
            nn.Conv2d(3, 64, 3, 1, 1),
            nn.BatchNorm2d(64),
            nn.ReLU(),
        )
        self.layer1 = self._make_layer(64, 64, 2)
        self.layer2 = self._make_layer(64, 128, 2, 2)
        self.layer3 = self._make_layer(128, 256, 2, 2)
        self.layer4 = self._make_layer(256, 512, 2, 2)
        self.avg_pool = nn.AvgPool2d(4)
        self.linear = nn.Linear(512, num_class)

    def _make_layer(self, in_channel, out_channel, bloch_num, stride=1):
        blocks = []
        blocks.append(ResidualBlock(in_channel, out_channel, stride))
        for i in range(1, bloch_num):
            blocks.append(ResidualBlock(out_channel, out_channel))
        return nn.Sequential(*blocks)

    def forward(self, x):
        x = self.conv1(x)
        x = self.layer1(x)
        x = self.layer2(x)
        x = self.layer3(x)
        x = self.layer4(x)
        x = self.avg_pool(x)

```



```
x = x.view(x.shape[0], -1)
return x
```

aug1

```
In [9]: trainloader, valloader = get_train_valid_loader('./data', batch_size=10,
augment=1, random_seed=34567,
                                                    valid_size=0.02, shuffle
= True, num_workers=2, pin_memory=True)


dataloaders = {'train': trainloader, 'val': valloader, 'test': testloader}
dataset_sizes = {'train': 49000, 'val': 1000, 'test': len(testset)}
```

Files already downloaded and verified

Files already downloaded and verified

```
In [39]: #Load Resnet18
model_ft = resnet()
#Finetune Final few layers to adjust for tiny imagenet input
# model_ft.avgpool = nn.AdaptiveAvgPool2d(1)
# model_ft.fc.out_features = 10
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
model_ft = model_ft.to(device)
#Loss Function
criterion = nn.CrossEntropyLoss()
# Observe that all parameters are being optimized
optimizer_ft = optim.SGD(model_ft.parameters(), lr=0.001, momentum=0.9)

#Train
train_model("resnet_aug1", model_ft, dataloaders, dataset_sizes, criteri
on, optimizer_ft, num_epochs=10)
```



```
accuracy
    training          (min:    0.470, max:    0.862, cur:
0.862)
    validation        (min:    0.595, max:    0.848, cur:
0.848)
log loss
    training          (min:    0.396, max:    1.462, cur:
0.396)
    validation        (min:    0.436, max:    1.061, cur:
0.436)
Train Loss: 0.3962 Acc: 0.8623
Val Loss: 0.4358 Acc: 0.8480
```

```
Training complete in 33m 52s
Best Validation Accuracy: 0.848, Epoch: 10
```

```
In [40]: #Test Resnet18
model_ft = resnet()
#Finetune Final few layers to adjust for tiny imagenet input
# model_ft.avgpool = nn.AdaptiveAvgPool2d(1)
# model_ft.fc.out_features = 10
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
model_ft.load_state_dict(torch.load('models/resnet_aug1/model_10_epoch.p
t')) #TODO
model_ft = model_ft.to(device)

#Loss Function
criterion = nn.CrossEntropyLoss()
# Observe that all parameters are being optimized
optimizer_ft = optim.SGD(model_ft.parameters(), lr=0.001, momentum=0.9)

#Test
test_model(model_ft, dataloaders, dataset_sizes, criterion, optimizer_ft
)
```

Iteration: 2500/2500, Loss: 3.2925541400909424...

Test Loss: 0.7696 Acc: 0.7391

Test complete in 0m 29s

aug2

```
In [10]: trainloader, valloader = get_train_valid_loader('./data', batch_size=10,
augment=2, random_seed=34567,
valid_size=0.02, shuffle
=True, num_workers=2, pin_memory=True)


dataloaders = {'train': trainloader, 'val': valloader, 'test': testload
er}
dataset_sizes = {'train': 49000, 'val': 1000, 'test': len(testset)}
```

Files already downloaded and verified

Files already downloaded and verified

```
In [42]: #Load Resnet18
model_ft = resnet()
#Finetune Final few layers to adjust for tiny imagenet input
# model_ft.avgpool = nn.AdaptiveAvgPool2d(1)
# model_ft.fc.out_features = 10
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
model_ft = model_ft.to(device)
#Loss Function
criterion = nn.CrossEntropyLoss()
# Observe that all parameters are being optimized
optimizer_ft = optim.SGD(model_ft.parameters(), lr=0.001, momentum=0.9)

#Train
train_model("resnet_aug2", model_ft, dataloaders, dataset_sizes, criteri
on, optimizer_ft, num_epochs=10)
```



```
accuracy
    training          (min:    0.408, max:    0.790, cur:
0.790)
    validation        (min:    0.560, max:    0.783, cur:
0.779)
log loss
    training          (min:    0.597, max:    1.612, cur:
0.597)
    validation        (min:    0.631, max:    1.207, cur:
0.631)
Train Loss: 0.5975 Acc: 0.7901
Val Loss: 0.6315 Acc: 0.7790
```

```
Training complete in 34m 33s
Best Validation Accuracy: 0.783, Epoch: 9
```

```
In [11]: #Test Resnet18
model_ft = resnet()
#Finetune Final few layers to adjust for tiny imagenet input
# model_ft.avgpool = nn.AdaptiveAvgPool2d(1)
# model_ft.fc.out_features = 10
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
model_ft.load_state_dict(torch.load('models/resnet_aug2/model_9_epoch.p
t')) #TODO
model_ft = model_ft.to(device)

#Loss Function
criterion = nn.CrossEntropyLoss()
# Observe that all parameters are being optimized
optimizer_ft = optim.SGD(model_ft.parameters(), lr=0.001, momentum=0.9)

#Test
test_model(model_ft, dataloaders, dataset_sizes, criterion, optimizer_ft
)
```

Iteration: 2500/2500, Loss: 2.306704521179199....

Test Loss: 0.9839 Acc: 0.6721

Test complete in 0m 34s

aug3

```
In [12]: trainloader, valloader = get_train_valid_loader('./data', batch_size=10,
augment=3, random_seed=34567,
valid_size=0.02, shuffle
=True, num_workers=2, pin_memory=True)


dataloaders = {'train': trainloader, 'val': valloader, 'test': testloade
r}
dataset_sizes = {'train': 49000, 'val': 1000, 'test': len(testset)}
```

Files already downloaded and verified

Files already downloaded and verified

```
In [47]: #Load Resnet18
model_ft = resnet()
#Finetune Final few layers to adjust for tiny imagenet input
# model_ft.avgpool = nn.AdaptiveAvgPool2d(1)
# model_ft.fc.out_features = 10
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
model_ft = model_ft.to(device)
#Loss Function
criterion = nn.CrossEntropyLoss()
# Observe that all parameters are being optimized
optimizer_ft = optim.SGD(model_ft.parameters(), lr=0.001, momentum=0.9)

#Train
train_model("resnet_aug3", model_ft, dataloaders, dataset_sizes, criteri
on, optimizer_ft, num_epochs=10)
```



```
accuracy
    training                (min:    0.378, max:    0.748, cur:
0.748)
    validation              (min:    0.486, max:    0.779, cur:
0.779)
log loss
    training                (min:    0.724, max:    1.700, cur:
0.724)
    validation              (min:    0.628, max:    1.476, cur:
0.628)
Train Loss: 0.7242 Acc: 0.7476
Val Loss: 0.6283 Acc: 0.7790
```

```
Training complete in 34m 47s
Best Validation Accuracy: 0.779, Epoch: 10
```

```
In [48]: #Test Resnet18
model_ft = resnet()
#Finetune Final few layers to adjust for tiny imagenet input
# model_ft.avgpool = nn.AdaptiveAvgPool2d(1)
# model_ft.fc.out_features = 10
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
model_ft.load_state_dict(torch.load('models/resnet_aug3/model_10_epoch.p
t')) #TODO
model_ft = model_ft.to(device)

#Loss Function
criterion = nn.CrossEntropyLoss()
# Observe that all parameters are being optimized
optimizer_ft = optim.SGD(model_ft.parameters(), lr=0.001, momentum=0.9)

#Test
test_model(model_ft, dataloaders, dataset_sizes, criterion, optimizer_ft
)
```

Iteration: 2500/2500, Loss: 4.481856822967529...

Test Loss: 1.5099 Acc: 0.5551

Test complete in 0m 28s

aug4

```
In [13]: trainloader, valloader = get_train_valid_loader('./data', batch_size=10,
augment=4, random_seed=34567,
valid_size=0.02, shuffle
=True, num_workers=2, pin_memory=True)


dataloaders = {'train': trainloader, 'val': valloader, 'test': testloade
r}
dataset_sizes = {'train': 49000, 'val': 1000, 'test': len(testset)}
```

Files already downloaded and verified

Files already downloaded and verified

```
In [50]: #Load Resnet18
model_ft = resnet()
#Finetune Final few layers to adjust for tiny imagenet input
# model_ft.avgpool = nn.AdaptiveAvgPool2d(1)
# model_ft.fc.out_features = 10
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
model_ft = model_ft.to(device)
#Loss Function
criterion = nn.CrossEntropyLoss()
# Observe that all parameters are being optimized
optimizer_ft = optim.SGD(model_ft.parameters(), lr=0.001, momentum=0.9)

#Train
train_model("resnet_aug4", model_ft, dataloaders, dataset_sizes, criteri
on, optimizer_ft, num_epochs=10)
```



```
accuracy
  training      (min:    0.328, max:    0.658, cur:
0.658)
  validation    (min:    0.431, max:    0.699, cur:
0.699)
log loss
  training      (min:    0.968, max:    1.815, cur:
0.968)
  validation    (min:    0.838, max:    1.543, cur:
0.838)
Train Loss: 0.9682 Acc: 0.6578
Val Loss: 0.8378 Acc: 0.6990
```

```
Training complete in 33m 38s
Best Validation Accuracy: 0.6990000000000001, Epoch: 10
```



```
In [51]: #Test Resnet18
model_ft = resnet()
#Finetune Final few layers to adjust for tiny imagenet input
# model_ft.avgpool = nn.AdaptiveAvgPool2d(1)
# model_ft.fc.out_features = 10
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
model_ft.load_state_dict(torch.load('models/resnet_aug4/model_10_epoch.p
t')) #TODO
model_ft = model_ft.to(device)

#Loss Function
criterion = nn.CrossEntropyLoss()
# Observe that all parameters are being optimized
optimizer_ft = optim.SGD(model_ft.parameters(), lr=0.001, momentum=0.9)

#Test
test_model(model_ft, dataloaders, dataset_sizes, criterion, optimizer_ft
)
```

Iteration: 2500/2500, Loss: 4.068412780761719...

Test Loss: 1.6114 Acc: 0.4739

Test complete in 0m 29s